

Findings of the MAGMaR 2026 Shared Task

Alexander Martin¹ Dengjia Zhang¹ Joel Brogan^{2*} Francis Ferraro³
Jeremy Gwinnup⁴ Reno Kriz⁵ Teng Long⁶
Kenton Murray⁵ Andrew Yates⁵ Xiang Xiang⁷

¹Johns Hopkins University ²OpenAI

³University of Maryland, Baltimore County ⁴Air Force Research Laboratory

⁵Human Language Technology Center of Excellence, Johns Hopkins University

⁶University of Amsterdam ⁷Huazhong University of Science and Technology

{amart233, kenton, rkriz1}@jhu.edu

Abstract

This overview paper presents the results of the shared task for the second workshop on Multimodal Augmented Generation via Multimodal Retrieval (MAGMaR). In this shared task participants submitted systems focused on either (i) video retrieval or (ii) grounded generation of articles given retrieved videos. Teams could submit to either task. For the retrieval task, we had 2 participating teams that submitted a total of 17 systems – all of which beat a baseline derived from the winner of last year’s shared task. On the generation side, we had 4 teams submit 16 systems. All teams had at least one system that received human annotation votes ranking at least 6 of the generated articles as being the best for that query.

1 Introduction

The second workshop on Multimodal Augmented Generation via Multimodal Retrieval (MAGMaR 2026) was located at ACL in San Diego and hosted a shared task focused on video retrieval and generation. Increasingly, information seeking needs require solutions that span modalities, yet much of the literature, and open shared tasks, continue to focus on one modality (i.e., text). The focus of this shared task is aimed at addressing this deficiency. The task builds upon the success of last year’s shared task of retrieving videos given a query by expanding the test collection of videos, and then introducing a new generation track:

- **Retrieval Track:** Systems provide a ranked list of videos in the collection ordered by relevance to the query.
- **Generation Track:** Systems produce a text article that answers the information need and grounds its content in the retrieved videos.

This shared task focuses on retrieving relevant videos and generating grounded articles that respond to information needs. Given a query describing a real-world current event, participating systems must identify pertinent videos from a large multilingual, multimodal collection and use that evidence to produce a coherent and informative written article. Our retrieval collection comprised over 110,000 multilingual, event-centric videos. The generation task had 19 queries with which systems needed to produce an article that met the information seeking needs of a specific persona. The summaries were judged by 3 human participants.

Overall, this was a very successful shared task garnering multiple submissions with all teams beating very strong baselines in both tracks. Teams made meaningful improvements and explored cutting-edge solutions to an open-problem. We are delighted with the results and suggest you read the system descriptions. That being said, here are some of the highlights across submissions:

Findings of the 2026 MAGMaR Shared Task

- Text-based reasoning dominated both tracks. Converting videos into summaries and captions rather than operating on video directly helps in both retrieval and generation.
- Complex, persona-constrained queries against a large video corpus do not perform well using only a single embedding. However, more expressive retrieval architectures can do substantially better.
- The best performing retrieval systems relied on broad first-stage recall with reasoning-based reranking and significantly outperforms all baselines.
- There is a clear disconnect between human preference and automatic metrics for scoring

*Subsequent authors listed alphabetically

generations. Systems that abstain when evidence is insufficient are rewarded by annotators yet penalized by the automatic metric as precision failures.

2 Data

The shared task is built on WIKIVIDEO25, which extends WIKIVIDEO (Martin et al., 2025) to address its limitations in query specificity, language coverage, video diversity, and annotation methodology. Figure 1 gives a general visualization, and Appendix A provides a detailed comparison with the original dataset.

Queries and personas. Each topic comes with a persona and a query (two per topic; see Figure 2). The persona constrains what counts as relevant evidence, so systems must reason about which parts of the videos matter for a specific professional need. Each query is written in one of two ways: *biased*, where the author had watched the videos, and *unbiased*, where the author had only seen video titles. Unbiased queries may ask for information absent from the collection, testing whether systems can acknowledge gaps rather than hallucinate.

Annotation and video diversity. Ground-truth articles are built bottom-up: annotators first write atomic claims grounded in each video’s audiovisual content, then combine them into higher-level inferences. This produces reports faithful to the videos and reduces the advantage of parametric knowledge. The video collection shifts toward raw and “diet raw” footage, videos captured by bystanders without professional narration, which force models to infer events from low-level audiovisual cues instead of relying on transcripts alone. WIKIVIDEO25 also adds two Mandarin-language topics and one mixed Nepali–English topic, so systems must process non-English audio and OCR.

MultiVENT 2.0. Last year’s shared task only had the retrieval track. We used the test collection of MultiVENT 2.0 (Kriz et al., 2025), a dataset of over 110,000 videos across 6 languages. This year, these videos were used as a distractor set which was combined with WIKIVIDEO25 to form the test collection. The videos in MultiVENT2.0 and WIKIVIDEO25 come from similar sources, video types, and distributions. Because MultiVENT 2.0 covers events only through 2023 and all WIKIVIDEO25 events occur in 2025, no distractor video can be a true positive for any query;

relevance judgments therefore apply exclusively to the WIKIVIDEO25 portion of the collection.

3 Retrieval Task

Task. The retrieval task requires systems to rank a corpus of over 110,000 multilingual, event-centric videos by relevance to a structured query. As illustrated in Figure 2, each query comes with a detailed persona and a complex information need. Systems must return a ranked list of the top 1000 videos and are evaluated by nDCG and Recall at cutoffs of 10, 20, and 100, measured using ir-measures (MacAvaney et al., 2022), over the combined WIKIVIDEO25 and MultiVENT 2.0 (Kriz et al., 2025) test corpora.

3.1 Systems

We provide participants with three baselines: OmniEmbed (Ma et al., 2025), an open-source dense retriever; OmniEmbed + RankVideo (Skow et al., 2026), OmniEmbed reranked by an open-source video-native reranker; and Wholembd-v3 (Mixedbread Team, 2026), a closed-source first-stage retrieval model.

We received 17 submissions¹ from two systems: C2F-RAG (Dai et al., 2026) and MARQUIS (Chakraborty et al., 2026). Both follow a similar pipeline: broad first-stage recall followed by reasoning-based reranking, with no task-specific training. C2F-RAG indexes text proxies of the video content (global summaries and keyframe captions) with BGE-M3 (Chen et al., 2025), then applies an LLM-based cognitive reranking agent that scores each candidate on logical alignment with the query persona. MARQUIS operates on the video content directly, decomposing each query into atomic sub-queries, retrieving independently over each with OmniEmbed, fusing the resulting ranked lists, and reranking with RankVideo.

3.2 Results

Table 1 summarizes the retrieval results.

First-stage recall vs. reasoning-based reranking.

The baselines span a wide performance range. OmniEmbed achieves an nDCG@10 of 0.17, which shows how difficult it is to match complex, persona-constrained queries against a large video corpus with a single embedding. Adding RankVideo as a reranker more than triples this to 0.54; even over a weak first stage, a reasoning-based second stage

¹Only showing the top 5 from MARQUIS



Figure 1: WIKIVIDEO25 Summary Statistics. “Language” (left): Videos are in English, Chinese, and Nepali. “Event Type & Topic” (center): are partitioned by number of videos per topic and color coded by event type. Event Type Key: **Orange**: Natural Disasters, **Blue**: Political Developments, **Green**: Science. “Video Type” (right): Nearly half of videos are raw and diet raw footage, captured by bystanders without professional editing or narration.

Query 1

PERSONA Statistician for North American Elections

QUERY Help me compile parliamentary and vote share statistics on the 2025 Canadian Federal elections. This should include how many seats each party won or lost, how many total seats each party now has, any information on total (popular) vote share available for the major parties and for any specific candidates for which it is available. Any demographic breakdowns of support for particular parties or candidates would also be helpful. For all statistics, please also include as much detail as possible on their source, as this helps me get a sense for how credible they are.

Query 2

PERSONA Policy Analyst at Elections Canada

QUERY I am looking to understand the key dynamics and outcomes of the 2025 Canadian Federal Election. Specifically, I am interested in how vote shares translated into seat outcomes under the first-past-the-post system, the regional patterns that enabled the Liberals to form government, and the factors contributing to the sharp decline in NDP support. Information on polling trends during the campaign, the role of leadership changes, and the influence of external political pressures—such as U.S. political rhetoric—would be particularly valuable.

Figure 2: Two example queries from WIKIVIDEO25 for the 2025 Canadian Federal Election event. Query 1 is *unbiased* and adopts a statistician persona seeking numerical data; Query 2 is *biased* and adopts a policy analyst persona focused on structural explanations. Both target English-language output.

makes a large difference. Wholembed-v3, a closed-source late-interaction model, reaches 0.72 without any reranking, so more expressive retrieval architectures can close much of the gap on their own. Both submitted systems pair broad first-stage recall with reasoning-based reranking and outperform all baselines.

Text proxies vs. video-native retrieval. C2F-RAG achieves the highest scores on every metric (nDCG@10 of 0.848) despite never operating on video content directly; it indexes text proxies (global summaries and keyframe captions) and reranks with an LLM-based cognitive agent. MARQUIS operates on video directly through OmniEmbed, yet its best configuration (MARQUIS-2) reaches only 0.759. Whether text proxies are a better retrieval substrate for complex queries than cur-

rent video embeddings, or whether any reasonable first stage would perform similarly when paired with a strong reranker, remains an open question.

Query decomposition and rank fusion. MARQUIS decomposes each query into atomic subqueries, retrieves independently over each, and fuses the resulting ranked lists before reranking. Starting from OmniEmbed, the weakest first-stage baseline (0.17), this strategy closes the gap: MARQUIS-2 reaches 0.759. The tight clustering of MARQUIS variants (0.746–0.759), however, suggests that the gains come mainly from query expansion and reranking, not from the choice of fusion strategy.

Several questions remain. Combining a stronger first-stage model like Wholembed-v3 with C2F-RAG’s cognitive reranking, or with a trained video-

Rank	System	nDCG@10	nDCG@20	nDCG@100	R@10	R@20	R@100
–	<i>OmniEmbed</i>	<i>0.166</i>	<i>0.186</i>	<i>0.245</i>	<i>0.096</i>	<i>0.158</i>	<i>0.297</i>
–	<i>OmniEmbed + RankVideo</i>	<i>0.542</i>	<i>0.534</i>	<i>0.546</i>	<i>0.423</i>	<i>0.462</i>	<i>0.494</i>
–	<i>Mixedbread</i>	<i>0.717</i>	<i>0.706</i>	<i>0.748</i>	<i>0.604</i>	<i>0.634</i>	<i>0.741</i>
6	MARQUIS-6	0.746	0.757	0.799	0.636	0.711	0.818
5	MARQUIS-8	0.747	0.758	0.800	0.636	0.711	0.818
4	MARQUIS-4	0.754	0.765	0.807	0.641	0.716	0.823
3	MARQUIS-14	0.757	0.768	0.810	0.650	0.725	0.832
2	MARQUIS-2	0.759	0.771	0.811	0.652	0.730	0.832
1	C2F-RAG	0.848	0.851	0.853	0.773	0.832	0.837

Table 1: MAGMaR 2026 — Retrieval track leaderboard (sorted by nDCG@10). Baselines shown in italics above the rule. Note that we received 17 submission and all of them beat a strong baseline based on the winner of last year’s shared task. As many submitted systems were slight variants of each other, we are only showing the top 6 systems from the leaderboard.

native reranker, would help separate the contributions of recall from reasoning. More broadly, the results point to a gap in multimodal embedding models and rerankers for complex retrieval over long-form video.

4 Generation Task

Task. The generation task requires systems to produce a cited, persona-consistent article grounded in a set of relevant videos. In the *oracle* setting, systems receive the ground-truth relevant videos for each query, isolating generation quality from upstream retrieval variance. Generated articles are evaluated both by human annotators and by MiRAGE (Martin et al., 2026), an automatic framework that measures factuality, information coverage, groundedness, and citation attribution. Each MiRAGE entailment judgment is produced by CLUE (Zhang et al., 2026), a calibrated multimodal uncertain-inference model. Full results are in Table 2.

4.1 Systems

We provide Collaborative Article Generation (CAG) (Martin et al., 2025) as a baseline with Qwen3.5 as the backbone (Qwen Team, 2026). CAG achieves a human score of 3.088 with strong precision on both information (0.76) and citation (0.62) metrics, but low recall.

Four teams submitted 16 predictions to the leaderboard: TRACE (Yan et al., 2026), CRAFT (Bhosale et al., 2026), C2F-RAG (Dai et al., 2026), and MARQUIS (Chakraborty et al., 2026). TRACE builds a structured text timeline for each video via OCR and object detection, then uses a text-only LLM to localize relevant moments before passing selected frames to a vision-language model for

claim generation. CRAFT generates atomic claims from query-conditioned keyframes and ASR, then iteratively verifies and repairs each claim through a hybrid critic loop before merging citations. C2F-RAG applies its retrieval pipeline’s prompt sculpting mechanism to constrain generation to the persona and enforce chunk-level temporal grounding. MARQUIS explores multiple generation strategies over a shared evidence extraction layer, including direct QA, clustering-based summarization, and an RLM controller that iteratively gathers and curates evidence before writing.

4.2 Results

The generation results (Table 2) show several patterns across the submissions

Human preference vs. automatic metrics.

The two top-ranked systems by human score, MARQUIS-iter-qa-base (3.833) and TRACE-04 (3.825), achieve relatively modest MiRAGE precision and recall compared to lower-ranked systems. Conversely, the CAG baseline and MARQUIS-Ginger achieve the highest InfoP (0.764 and 0.776) and CiteP (0.617 and 0.643), yet rank 6th and below by human score. The gap comes mainly from how QA-based systems handle unanswerable sub-questions: when the video evidence is insufficient, these systems refuse to answer rather than hallucinate, and human annotators reward this. MiRAGE, however, scores refusals as missing information and penalizes recall. So the systems human annotators prefer most, precisely because they acknowledge gaps rather than fabricate claims, are the ones automatic evaluation penalizes most.

QA-based generation. The MARQUIS QA variants dominate the top of the human leaderboard,

Rank	System	Human	Best Votes	Best %	InfoP	InfoR	CiteP	CiteR
–	<i>CAG (baseline)</i>	<i>3.088</i>	<i>1</i>	<i>1.8%</i>	<i>0.764</i>	<i>0.410</i>	<i>0.617</i>	<i>0.228</i>
16	CRAFT-01	2.246	0	0.0%	0.695	0.389	0.065	0.000
15	CRAFT-02	2.281	0	0.0%	0.627	0.304	0.530	0.198
14	C2F-RAG-01	2.456	6	10.5%	0.557	0.466	0.452	0.349
13	C2F-RAG-02	2.526	2	3.5%	0.584	0.450	0.479	0.347
12	CRAFT-03	2.542	0	0.0%	0.609	0.304	0.509	0.204
11	TRACE-02	2.579	0	0.0%	0.621	0.550	0.536	0.500
10	TRACE-01	2.579	0	0.0%	0.599	0.564	0.515	0.498
9	MARQUIS-Bullet	2.667	0	0.0%	0.711	0.394	0.604	0.237
8	TRACE-03	2.702	2	3.5%	0.599	0.553	0.507	0.503
7	MARQUIS-ss-qa-base	3.070	6	10.5%	0.331	0.306	0.277	0.281
6	MARQUIS-Ginger	3.123	6	10.5%	0.776	0.404	0.643	0.226
5	MARQUIS-RLM	3.298	3	5.3%	0.708	0.385	0.592	0.272
4	MARQUIS-ss-qa-ginger	3.421	10	17.5%	0.544	0.324	0.326	0.238
3	MARQUIS-iter-qa-ginger	3.694	5	8.8%	0.345	0.290	0.257	0.226
2	TRACE-04	3.825	8	14.0%	0.640	0.483	0.498	0.405
1	MARQUIS-iter-qa-base	3.833	8	14.0%	0.347	0.313	0.268	0.258

Table 2: MAGMaR 2026 — Oracle generation track leaderboard (sorted by human score). Baseline shown in italics above the rule. 4 teams submitted systems. All teams had at least one system which earned best votes.

holding three of the top four positions. These systems decompose the query into sub-questions, answer each from the video evidence, and synthesize the results into an article. Annotators find these outputs more complete and easier to evaluate, likely because the question-answer structure gives the article a natural organization. The gap between iterative and single-shot QA variants (3.833 vs. 3.070) suggests that iterating over the evidence, asking follow-up questions and refining answers, produces noticeably better articles.

Evidence verification. Systems with explicit claim verification show stronger citation precision. CRAFT-03 achieves a CiteP of 0.509 despite ranking 12th overall, and MARQUIS-Ginger reaches 0.643 with its CLUE-based calibration filtering. The TRACE submissions achieve the most balanced precision-recall profiles across both information and citation metrics: TRACE-04 reaches 0.640 InfoP / 0.483 InfoR and 0.498 CiteP / 0.405 CiteR, the strongest combined recall among top-ranked systems.

The shared task results point to several open problems. The clearest is the disconnect between human preference and automatic metrics: systems that abstain when evidence is insufficient are rewarded by annotators but penalized by MiRAGE, which treats all omissions as precision failures. Evaluation frameworks that account for appropriate refusal and persona-relevant (not merely factual)

coverage are a clear next step. On the modeling side, no submitted system uses task-specific fine-tuning; the strong performance of QA-based approaches suggests that models trained for grounded multi-video question answering with persona conditioning could improve results. Finally, all current systems convert videos to text proxies before reasoning, discarding visual information that captions cannot capture. Architectures for direct multi-video visual reasoning are needed to close this gap.

5 Conclusion

The shared task at the second workshop on Multimodal Augmented Generation via Multimodal Retrieval (MAGMaR 2026) attracted two retrieval teams and four generation teams. Text-based reasoning dominates both tracks: C2F-RAG achieves the highest retrieval scores by indexing summaries and captions rather than operating on video directly, and every generation system converts videos to text before reasoning. Given this shared substrate, inference-time reasoning proves more important than the choice of first-stage model. Reranking transforms OmniEmbed from the weakest retrieval baseline into a competitive first stage, and iterative QA generation produces noticeably better articles than single-shot variants. On the generation side, human preference and automatic metrics also diverge: systems that appropriately abstain when evidence is insufficient are rewarded by annotators but penalized by MiRAGE, which treats all omissions

as failures.

These results point to several directions for future work. Evaluation frameworks need to account for appropriate refusal and persona-relevant coverage beyond factual completeness. No submitted system uses task-specific fine-tuning in either track, and the strong performance of QA-based generation suggests that models trained for grounded, persona-conditioned multi-video question answering could improve further. All current systems also discard visual information by converting videos to text before reasoning; architectures for direct multi-video visual reasoning remain unexplored.

References

- Mahesh Bhosale, Abdul Wasi, Vishvesh Trivedi, Pengyu Yan, Akhil Gorugantu, and David Doermann. 2026. [Craft: Critic-refined adaptive key-frame targeting for multimodal video question answering](#). *Preprint*, arXiv:2605.19075.
- Debashish Chakraborty, Dengjia Zhang, Jialiang Jin, Hanting Liu, Katherine Guerriero, Hanxiang Qin, Tyler Skow, Alexander Martin, Reno Kriz, and Benjamin Van Durme. 2026. [Marquis: A three-stage pipeline for video retrieval-augmented generation](#). *Preprint*, arXiv:2605.17640.
- Jianlv Chen, Shitao Xiao, Peitian Zhang, Kun Luo, Defu Lian, and Zheng Liu. 2025. [M3-embedding: Multi-linguality, multi-functionality, multi-granularity text embeddings through self-knowledge distillation](#). *Preprint*, arXiv:2402.03216.
- Jiixin Dai, Zehang Wei, Jiamin Yan, and Xiang Xiang. 2026. [Decoupling semantics and logic: A training-free coarse-to-fine pipeline for video retrieval-augmented generation](#). *Arxiv preprint*, 2606.07924.
- Reno Kriz, Kate Sanders, David Etter, Kenton Murray, Cameron Carpenter, Kelly Van Ochten, Hannah Recknor, Jimena Guallar-Blasco, Alexander Martin, Ronald Colaanni, Nolan King, Eugene Yang, and Benjamin Van Durme. 2025. [Multivent 2.0: A massive multilingual benchmark for event-centric video retrieval](#). *Preprint*, arXiv:2410.11619.
- Dawn Lawrie, Sean MacAvaney, James Mayfield, Paul McNamee, Douglas W. Oard, Luca Soldaini, and Eugene Yang. 2025. [Overview of the trec 2024 neuclir track](#). *Preprint*, arXiv:2509.14355.
- Xueguang Ma, Luyu Gao, Shengyao Zhuang, Jiaqi Samantha Zhan, Jamie Callan, and Jimmy Lin. 2025. [Tevatron 2.0: Unified document retrieval toolkit across scale, language, and modality](#). *Preprint*, arXiv:2505.02466.
- Sean MacAvaney, Craig Macdonald, and Iadh Ounis. 2022. [Streamlining evaluation with ir-measures](#). In *Advances in Information Retrieval - 44th European Conference on IR Research, ECIR 2022, Stavanger, Norway, April 10-14, 2022, Proceedings, Part II*, volume 13186 of *Lecture Notes in Computer Science*, pages 305–310. Springer.
- Alexander Martin, Reno Kriz, William Gantt Walden, Kate Sanders, Hannah Recknor, Eugene Yang, Francis Ferraro, and Benjamin Van Durme. 2025. [Wikivideo: Article generation from multiple videos](#). *Preprint*, arXiv:2504.00939.
- Alexander Martin, William Walden, Reno Kriz, Dengjia Zhang, Kate Sanders, Eugene Yang, Chihsheng Jin, and Benjamin Van Durme. 2026. [Seeing through the mirage: Evaluating multimodal retrieval augmented generation](#). *Preprint*, arXiv:2510.24870.
- Mixedbread Team. 2026. [Beyond the limit: Introduce mixedbread wholembd v3](#). Blog post.
- Qwen Team. 2026. [Qwen3.5: Towards native multi-modal agents](#).
- Tyler Skow, Alexander Martin, Benjamin Van Durme, Rama Chellappa, and Reno Kriz. 2026. [Rankvideo: Reasoning reranking for text-to-video retrieval](#). *Preprint*, arXiv:2602.02444.
- Pengyu Yan, Akhil Gorugantu, Mahesh Bhosale, Abdul Wasi, Vishvesh Trivedi, and David Doermann. 2026. [Trace: Evidence grounding-guided multi-video event understanding and claim generation](#). *Preprint*, arXiv:2605.16740.
- Dengjia Zhang, Alexander Martin, William Jurayj, Kenton Murray, Benjamin Van Durme, and Reno Kriz. 2026. [Unified multimodal uncertain inference](#). *Preprint*, arXiv:2604.08701.

A WIKIVIDEO25 Dataset

WIKIVIDEO25 builds on WIKIVIDEO (Martin et al., 2025), which contains 57 English topics from 2015–2023, with ~ 7 videos per topic. It has several limitations: it is English-only, its articles are written top-down from preexisting Wikipedia lead sections, its videos are mostly professional news broadcasts, and its queries are underspecified (e.g., “tell me about [event name]”) and do not reflect realistic information-seeking behavior (Lawrie et al., 2025). Because the topics are well-known events, the relevant information often already exists in pre-trained model weights. In some cases, simply quoting the Wikipedia lead section outperforms human-written responses. WIKIVIDEO25 targets each of these problems.

Queries and Personas. Each topic comes with a persona and a query (two per topic; see Figure 2). The persona is a detailed background for the query

provider. It constrains what counts as relevant evidence, so systems must reason about which parts of the videos matter for a specific professional need, not just surface everything about the event. Each query is a complex information need written in one of two ways: *biased*, where the author had watched the videos and knows what the collection contains, and *unbiased*, where the author had only seen video titles and writes a more natural request without that knowledge. Unbiased queries may ask for information absent from the collection, so they also test whether systems can acknowledge gaps in the evidence instead of hallucinating an answer.

Video-Centric Reports. WIKIVIDEO’s articles were written top-down: annotators began with the Wikipedia lead section for each event and removed claims not supported by the videos. The ground-truth articles were therefore anchored to an external text source, not to the video content itself, and in some cases quoting the Wikipedia lead outperformed human-written responses. WIKIVIDEO25 reverses this. Annotators first write atomic claims grounded in each video’s audiovisual content, then combine them into higher-level inferences that form the final article. This bottom-up annotation produces reports more faithful to the video content and reduces the advantage of parametric knowledge, since the ground truth is no longer derivable from a well-known Wikipedia article.

Knowledge cutoff. All events in WIKIVIDEO25 take place in 2025, after the training cutoffs of the models used in submitted systems. Because none of the event-level facts exist in these models’ parametric knowledge, systems cannot shortcut the task by recalling memorized information and must ground their outputs in the retrieved video evidence. This resolves the quoting issue noted above from WIKIVIDEO, whose topics were well-known events with extensive coverage in pretraining data. Using events beyond any model’s knowledge cutoff removes this issue. It also means the Multi-VENT 2.0 distractor set, which covers events only through 2023, contains no true positives for any WIKIVIDEO25 query. The result is a cleaner evaluation of both retrieval and generation: retrieval systems must identify relevant videos rather than match familiar event descriptions, and generation systems must synthesize articles from audiovisual content rather than parametric knowledge.

Raw Videos. The original dataset is mostly professional news broadcasts, where a third-party narrator describes the event to the viewer. For these, an audio transcript alone is often enough to produce a reasonable article. WIKIVIDEO25 shifts toward raw and “diet raw” footage: videos captured by participants or bystanders without professional narration or post-production editing. These videos are more ambiguous and force models to infer what is happening from low-level audiovisual cues. There is no narrator explaining the scene.

Multilingual Videos. WIKIVIDEO25 adds two topics with Mandarin-language videos and one mixed Nepali–English topic. Systems must now process non-English audio and OCR at both retrieval and generation. In the mixed-language case, they must also reason across videos in different languages about the same event.